

# Trees

Muchang Bahng

Spring 2025

## Contents

<b>1</b>	<b>Decision Trees</b>	<b>3</b>
1.1	Classification Trees . . . . .	3
1.2	Regression Trees . . . . .	5
1.3	Model Space . . . . .	5
<b>2</b>	<b>Greedy Optimization</b>	<b>6</b>
2.1	Surrogate Loss . . . . .	8
2.2	CART . . . . .	8
2.3	ID3 . . . . .	8
2.4	c4.5 . . . . .	8
<b>3</b>	<b>Regularization</b>	<b>9</b>
3.1	Pruning . . . . .	10
3.2	Splitting . . . . .	10
<b>4</b>	<b>Improved Optimization</b>	<b>11</b>
4.1	GODST . . . . .	11
<b>5</b>	<b>Soft Decision Trees</b>	<b>12</b>
5.1	Soft Splitting . . . . .	12
5.2	Neural Decision Trees . . . . .	12
	<b>References</b>	<b>13</b>

Say that you were looking at a picture and were told to classify it as a bird or a dog. You would probably look for some features and have the following thought process: if it has wings, then it is a bird, and if not, then it is a dog.

Now let's try a slightly harder classification. We have four classes consisting of two species of dogs (golden retriever, husky) and two species of birds (pigeon, hummingbird). Then you might work something like this.

1. If it has wings, and
  - (a) it has a long beak, then it is a hummingbird.
  - (b) it doesn't have a long beak, then it is a pigeon.
2. If it doesn't have wings, and
  - (a) its fur color is yellow, then it is a golden retriever.
  - (b) its fur color is not yellow, then it is a husky.

Decision trees attempt to model this method of thinking by using a tree structure, and hopefully this example should convince you that this type of model is worth studying. It is a discriminative model that learns to classify data by first identifying the relevant feature to look at (e.g. wings, beak length, fur color) and then deciding how to split it.

Surprisingly, the origin of tree models is not clear, though there have been some papers as early as 1959 that mentions a decision tree-like structure.<sup>1</sup> I personally would have thought it to be older given the simplicity of the idea.

---

<sup>1</sup>See <https://stats.stackexchange.com/questions/257537/who-invented-the-decision-tree>.

# 1 Decision Trees

In here, we define the decision tree model. It is most natural for classification, but there are variants of it for regression.

Many discriminative models can be written in a clean formula (e.g.  $y = w^T x + \epsilon$  for linear regression, and even  $y = \prod_i (\sigma_i \circ A_i)(x)$  for MLPs). However, we cannot find such a parametric form for a tree, which is why they are nonparametric models. In full generality, all we can say is that they have a general tree structure, and there are many variants.

## 1.1 Classification Trees

### Definition 1.1 (Classification Trees)

A **decision/classification tree** is a nonparametric discriminative model  $f$  that creates some sort of tree representing a set of decisions on an input  $x$  to predict a label  $y$ .

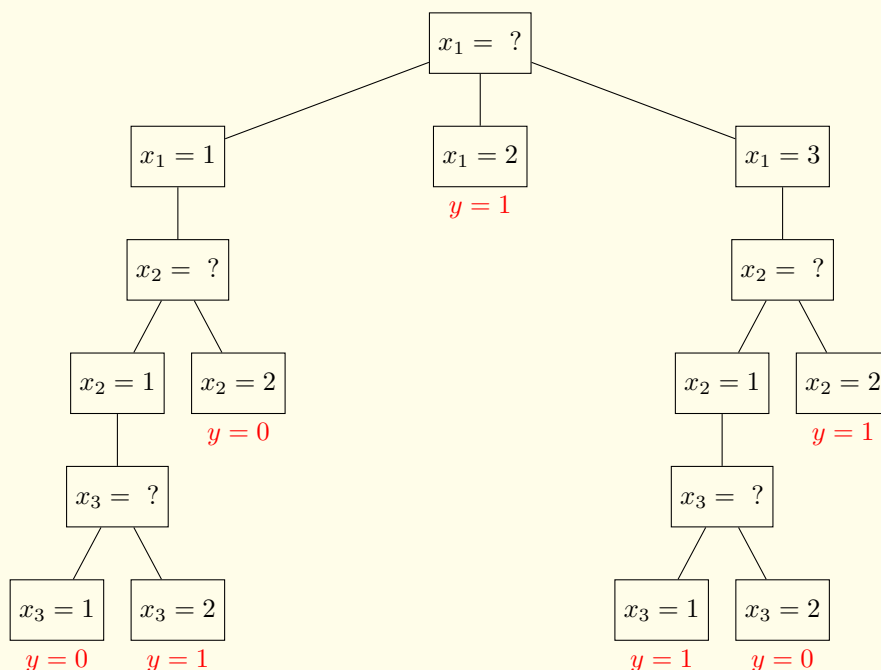


Figure 1: An example of a decision tree that splits at  $x_1$  first, then  $x_2$ , and finally  $x_3$ . Note that you can still split on  $x_2$  if  $x_1 = 1$  and  $x_3$  if  $x_1 = 3$ .

The decision tree tries to take advantage of some nontrivial covariance between  $X$  and  $Y$  by constructing nested partitions of the dataset  $\mathcal{D}$ , and within a partition, it predicts the label that comprises the majority.

### Example 1.1 (Restaurant Dataset)

Consider the following dataset, where we consider a binary classification system with discrete covariates.

	OthOptions	Weekend	WaitArea	Plans	Price	Precip	Restaur	Wait	Crowded	Stay?
$x_1$	Yes	No	No	Yes	\$\$\$	No	Mateo	0-5	some	Yes
$x_2$	Yes	No	No	Yes	\$	No	Juju	16-30	full	No
$x_3$	No	No	Yes	No	\$	No	Pizza	0-5	some	Yes
$x_4$	Yes	Yes	No	Yes	\$	No	Juju	6-15	full	Yes
$x_5$	Yes	Yes	No	No	\$\$\$	No	Mateo	30+	full	No
$x_6$	No	No	Yes	Yes	\$\$	Yes	BlueCorn	0-5	some	Yes
$x_7$	No	No	Yes	No	\$	Yes	Pizza	0-5	none	No
$x_8$	No	No	No	Yes	\$\$	Yes	Juju	0-5	some	Yes
$x_9$	No	Yes	Yes	No	\$	Yes	Pizza	30+	full	No
$x_{10}$	Yes	Yes	Yes	Yes	\$\$\$	No	BlueCorn	6-15	full	No
$x_{11}$	No	No	No	No	\$	No	Juju	0-5	none	No
$x_{12}$	Yes	Yes	Yes	Yes	\$	No	Pizza	16-30	full	Yes

Table 1: Dataset of whether to go to a restaurant for a date depending on certain factors.

Let us denote  $\mathcal{D}$  as the dataset, and say that  $F_1, \dots, F_d$  were the features. This is a binary classification problem, and we can count that there are 6 positives and 6 negative labels.

Note that this model is extremely flexible in that we can have different properties of these trees. We will introduce them as we go.

### Definition 1.2 (Binary Decision Tree)

A **binary decision tree** only allows the tree to split into two nodes.

Note that in density estimation or linear regression, we can derive the risk by first deriving the likelihood of the data, and then taking the negative logarithm of it to get our loss function which allows us to define our risk. In a decision tree, we have a *non-probabilistic* discriminative model, so there is no concept of likelihood. Therefore, we cannot use a pdf to define the loss. Fortunately, we can use the straightforward misclassification risk.

### Theorem 1.1 (Expected and Empirical Risk of Decision Trees)

Given a classification tree  $f$ , the misclassification risk over the true data generating distribution  $p(x, y)$ , along with its empirical risk over a dataset  $\mathcal{D} = (x^{(i)}, y^{(i)})_{i=1}^n$ , is

$$\mathbb{R}(f) = \mathbb{E}_{x,y} [\mathbb{1}(y \neq f(x))] = \int \mathbb{1}(y \neq f(x)) dx dy \quad (1)$$

$$\hat{R}(f) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}(y^{(i)} \neq f(x^{(i)})) \quad (2)$$

where  $\mathbb{1}(p)$  is an indicator function that equals 1 if  $p$  is true and 0 if false. Note that  $\hat{R}(f)$  is simply 1 minus the accuracy.

This is all nice in theory, but how are we supposed to optimize this in practice? First, the misclassification loss is not differentiable—and even worse—the gradient is 0 almost everywhere! This can be solved by using a surrogate loss function. Even worse,  $f$  is not parameteric, so we can't even gradients at all (with respect to what parameter?)! Another solution is to try and create a very specific tree model—making it parameteric—and then using a surrogate loss to learn. In fact this is what some people do, but for now we keep things simple. Therefore, we must use a gradient-free rule to optimize a decision tree.

## 1.2 Regression Trees

## 1.3 Model Space

The fact that trees are nonparametric means that we have extreme flexibility in designing our tree. However, this comes with the big risk of having too big of a model space to optimize over. This overcomplexity is one of the big challenges in trees.

For example, suppose that there are  $d$  covariates (independent variables, features)  $x_1, \dots, x_d$  all binary valued. We can design a decision tree that splits on  $x_1$ , then on  $x_2$ , then on  $x_1$ , then on  $x_2$ , and so on. This becomes unbounded and our model space a discrete infinite space, which is a bad combination since we don't have gradients to optimize over a continuum. We can try and handle this in two ways.

### Example 1.2 (Splitting on Same Variable Multiple Times)

Splitting on covariate  $x_1$  infinitely many times seems pretty unrealistic, so we should limit it in some way.

1. *Covariate can be split a maximum of once for each path from root to leaf.* This is a
2. *Covariate can be split maximum of  $k$  times.*

### Example 1.3 (Depth of Tree)

### Example 1.4 (Max Subnodes per Split)

## 2 Greedy Optimization

Optimizing a decision tree is NP-complete [HR76].

Let us take a decision tree  $h$  and model the accuracy of it as a random variable:  $1_{\{Y=h_0(X)\}} \sim \text{Bernoulli}(p)$ , where  $p$  is the accuracy. A higher accuracy of  $h$  corresponds to a lower entropy, and so the entropy of the random variable is also a relevant indicator.

$$H(1_{\{Y=h_0(X)\}}) = p \log p + (1 - p) \log(1 - p) \quad (3)$$

Therefore, when we are building a tree, we want to choose the feature  $x_i$  to split based on how much it lowers the entropy of the decision tree.

To set this up, let us take our dataset  $\mathcal{D}$  and set  $X_i$  as the random variable representing the distribution (a multinomial) of the  $x_i^{(j)}$ 's, and  $Y$  as the same for the  $y^{(j)}$ 's. This is our maximum likelihood approximation for the marginalized distribution of the joint measure  $X \times Y = X_1 \times \dots \times X_D \times Y$ .

Given a single node, we are simply going to label every point to be whatever the majority class is in  $\mathcal{D}$ . Therefore, we start off with the entropy of our trivial tree  $H(Y)$ . Then, we want to see which one of the  $X_d$  features to split on, and so we can compute the conditional entropy  $H(Y, X_d)$  to get the information gain  $I(Y; X_d) = H(Y) - H(Y | X_d)$  for all  $d = 1, \dots, D$ . We want to find a feature  $X_d$  that maximize this information gain, i.e. decreases the entropy as much as possible (a greedy algorithm), and we find the next best feature (with or without replacement), so that we have a decreasing sequence.

$$H(X) \geq H(X; Y) \geq H(X; Y, Z) \geq H(X; Y, Z, W) \geq \dots \geq 0$$

### Example 2.1 (Crowded Restaurants)

Continuing the example above, since there are 6 labels of 0 and 1 each, we can model this  $Y \sim \text{Bernoulli}(0.5)$  random variable, with entropy

$$H(Y) = \mathbb{E}[-\log_2 p(Y)] = \frac{1}{2} \left( -\log_2 \frac{1}{2} \right) + \frac{1}{2} \left( -\log_2 \frac{1}{2} \right) = 1 \quad (4)$$

Now what would happen if we had branched according to how crowded it was,  $X_{\text{crowded}}$ . Then, our decision tree would split into 3 sections:




Figure 2: Visual of decision tree splitting according to how crowded it is.

In this case, we can define the multinomial distribution  $X_{\text{crowded}}$  representing the proportion of the data that is crowded in a specific level. That is,  $X_{\text{crowded}} \sim \text{Multinomial}(\frac{2}{12}, \frac{4}{12}, \frac{6}{12})$ , with

$$\mathbb{P}(X_{\text{crowded}} = x) = \begin{cases} 2/12 & \text{if } x = \text{none} \\ 4/12 & \text{if } x = \text{some} \\ 6/12 & \text{if } x = \text{full} \end{cases} \quad (5)$$

Therefore, we can now compute the conditional entropy of this new decision tree conditioned on how crowded the store is

$$H(Y | X_{\text{crowded}}) = \sum_x \mathbb{P}(X_{\text{crowded}} = x) H(Y | X_{\text{crowded}} = x) \quad (6)$$

$$= \frac{2}{12} H(\text{Bern}(1)) + \frac{4}{12} H(\text{Bern}(0)) + \frac{6}{12} H(\text{Bern}(1/3)) = 0.459 \quad (7)$$

$$I(Y; X_{\text{crowded}}) = 0.541 \quad (8)$$

We would do this for all the features and greedily choose the feature that maximizes our information gain.

### Example 2.2 (Ferrari F1 Race)

The Ferrari F1 team hired you as a new analyst! You were given the following table of the past race history of the team. You were asked to use information gain to build a decision tree to predict race wins. First, you will need to figure out which feature to split first.

Rain	Good Strategy	Qualifying	Win Race
1	0	0	0
1	0	0	0
1	0	1	0
0	0	1	1
0	0	0	0
0	1	1	1
1	0	1	0
0	1	0	1
0	0	1	1
0	0	1	1

Let  $X \sim \text{Bernoulli}(1/2)$  be the distribution of whether a car wins a race over the data. Then its entropy is

$$H(X) = \mathbb{E}[-\log_2 p(x)] = \frac{1}{2}(-\log_2 \frac{1}{2}) + \frac{1}{2}(-\log_2 \frac{1}{2}) = 1 \quad (9)$$

Let  $R \sim \text{Bernoulli}(4/10)$ ,  $G \sim \text{Bernoulli}(2/10)$ ,  $Q \sim \text{Bernoulli}(6/10)$  be the distribution of the features rain, good strategy, and qualifying over the data, respectively. Then, the conditional entropy of  $X$  conditioned on each of these random variables is

$$\begin{aligned} H(X | R) &= \mathbb{P}(R = 1) H(X | R = 1) + \mathbb{P}(R = 0) H(X | R = 0) \\ &= \frac{4}{10} \cdot -(1 \cdot \log_2 1 + 0 \cdot \log_2 0) + \frac{6}{10} \cdot -\left(\frac{1}{6} \cdot \log_2 \frac{1}{6} + \frac{5}{6} \cdot \log_2 \frac{5}{6}\right) \approx 0.390 \end{aligned}$$

$$\begin{aligned} H(X | G) &= \mathbb{P}(G = 1) H(X | G = 1) + \mathbb{P}(G = 0) H(X | G = 0) \\ &= \frac{2}{10} \cdot -(1 \cdot \log_2 1 + 0 \cdot \log_2 0) + \frac{8}{10} \cdot -\left(\frac{3}{8} \cdot \log_2 \frac{3}{8} + \frac{5}{8} \log_2 \frac{5}{8}\right) \approx 0.763 \end{aligned}$$

$$\begin{aligned} H(X | Q) &= \mathbb{P}(Q = 1) H(X | Q = 1) + \mathbb{P}(Q = 0) H(X | Q = 0) \\ &= \frac{6}{10} \cdot -\left(\frac{4}{6} \cdot \log_2 \frac{4}{6} + \frac{2}{6} \cdot \log_2 \frac{2}{6}\right) + \frac{4}{10} \cdot -\left(\frac{1}{4} \log_2 \frac{1}{4} + \frac{3}{4} \log_2 \frac{3}{4}\right) \approx 0.875 \end{aligned}$$

Therefore, the information gain are

$$I(X; R) = 1 - 0.390 = 0.610$$

$$I(X; G) = 1 - 0.763 = 0.237$$

$$I(X; Q) = 1 - 0.875 = 0.125$$

And so I would split on  $R$ , the rain, which gives the biggest information gain.

Finally, we can use the Gini index of  $X \sim \text{Bernoulli}(p)$ , defined

$$G(X) = 2p(1 - p) \quad (10)$$

### Example 2.3 (Ferrari Example Continued)

We do the same as the Ferrari example above but now with the Gini reduction. Let  $X \sim \text{Bernoulli}(1/2)$  be the distribution of whether a car wins a race over the data. Then its Gini index, which I will label with  $\mathcal{G}$ , is

$$\mathcal{G}(X) = 2 \cdot \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{2}$$

Let  $R \sim \text{Bernoulli}(4/10)$ ,  $G \sim \text{Bernoulli}(2/10)$ ,  $Q \sim \text{Bernoulli}(6/10)$  be the distribution of the features rain, good strategy, and qualifying over the data, respectively. Then we compute the conditional expectation

$$\begin{aligned} \mathbb{E}[\mathcal{G}(X \mid R)] &= \mathbb{P}(R = 1) \mathcal{G}(X \mid R = 1) + \mathbb{P}(R = 0) \mathcal{G}(X \mid R = 0) \\ &= \frac{4}{10} \left[ 2 \cdot \frac{4}{4} \cdot \frac{0}{4} \right] + \frac{6}{10} \left[ 2 \cdot \frac{1}{6} \cdot \frac{5}{6} \right] \approx 0.167 \\ \mathbb{E}[\mathcal{G}(X \mid G)] &= \mathbb{P}(G = 1) \mathcal{G}(X \mid G = 1) + \mathbb{P}(G = 0) \mathcal{G}(X \mid G = 0) \\ &= \frac{2}{10} \left[ 2 \cdot \frac{2}{2} \cdot \frac{0}{2} \right] + \frac{8}{10} \left[ 2 \cdot \frac{3}{8} \cdot \frac{5}{8} \right] \approx 0.375 \\ \mathbb{E}[\mathcal{G}(X \mid Q)] &= \mathbb{P}(Q = 1) \mathcal{G}(X \mid Q = 1) + \mathbb{P}(Q = 0) \mathcal{G}(X \mid Q = 0) \\ &= \frac{6}{10} \left[ 2 \cdot \frac{4}{6} \cdot \frac{2}{6} \right] + \frac{4}{10} \left[ 2 \cdot \frac{1}{4} \cdot \frac{3}{4} \right] \approx 0.417 \end{aligned}$$

Therefore, the Gini reduction, which I'll denote as  $I_{\mathcal{G}}$ , is

$$\begin{aligned} I_{\mathcal{G}}(X; R) &= 0.5 - 0.167 = 0.333 \\ I_{\mathcal{G}}(X; G) &= 0.5 - 0.375 = 0.125 \\ I_{\mathcal{G}}(X; Q) &= 0.5 - 0.417 = 0.083 \end{aligned}$$

Since branching across the feature  $R$ , the rain, gives the biggest Gini reduction, we want to split on the rain feature first.

## 2.1 Surrogate Loss

## 2.2 CART

## 2.3 ID3

## 2.4 c4.5



### 3 Regularization

Given a dataset with  $D$  binary features, let  $g(H, D)$  be the number of binary trees with depth at most  $H$  (including root node), with the restriction that the trees may not split on some variable multiple times within a path to a leaf node. Then,  $g$  can be defined recursively.

1. First, if  $H = 1$ , then  $g(H, D) = 1$  always since we are just creating the trivial binary tree of one node.
2. If  $D = 0$ , then there are no features to split on and therefore we just have the single node  $g(H, D) = 1$ .
3. If  $H > 1$  and  $D > 0$ , then say that we start with a node. We can either make this a leaf node by not performing any splitting at all, or split on one of the  $D$  variables. Then for each of the 2 nodes created on the split, we are now working with  $D - 1$  features and a maximum height of  $H - 1$  for each of the subtrees generated from the 2 nodes.

All this can be expressed as

$$g(H, D) = \begin{cases} 1 + D [g(H - 1, D - 1)]^2 & \text{if } H > 1, D > 0 \\ 1 & \text{if } H = 1 \text{ or } D = 0 \end{cases}$$

which is extremely large (in fact, NP hard). Therefore, some tricks like regularization must be implemented to limit our search space.

By defining the complexity of our decision tree  $\Omega(h)$  as the number of nodes within the tree, we can modify our objective function to

$$L(h; \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N 1_{\{y^{(i)} \neq h(x^{(i)})\}} + \lambda \Omega(h)$$

We can impose this constraint directly on the training algorithm, or we can calculate the regularized loss after the tree has been constructed, which is a method called **tree pruning**.

Given a large enough  $\lambda$ , we can in fact greatly reduce our search space by not considering any trees further than a certain point.

#### Theorem 3.1 ()

We describe a tree as a set of leaves, where leaf  $k$  is a tuple containing the logical preposition satisfied by the path to leaf  $k$ , denoted  $p_k$ , and the class label predicted by the leaf, denoted  $\hat{y}_k$ . For a dataset with  $d$  binary features,  $p_k : \{0, 1\}^d \rightarrow \{0, 1\}$  is a function that returns 1 if a sample  $x_i$  satisfies the preposition, and 0 otherwise. That is, leaf  $k$  is  $(p_k, \hat{y}_k)$ , and a tree  $f$  with  $K$  leaves is described as a set  $f = \{(p_1, \hat{y}_1), \dots, (p_K, \hat{y}_K)\}$ . Assume that the label predicted by  $\hat{y}_k$  is always the label for the majority of samples satisfying  $p_k$ . Finally, let  $m_k = \sum_{i=1}^n p_k(x_i)$  denote the number of training samples “captured” by leaf  $k$ .

Given a (potentially optimal) tree

$$f = \{(p_1, \hat{y}_1), \dots, (p_\kappa, \hat{y}_\kappa), \dots, (p_K, \hat{y}_K)\},$$

the tree  $f' = \{(p_1, \hat{y}_1), \dots, (p_{\kappa_1}, \hat{y}_{\kappa_1}), (p_{\kappa_2}, \hat{y}_{\kappa_2}), \dots, (p_K, \hat{y}_K)\}$  produced by splitting leaf  $(p_\kappa, \hat{y}_\kappa)$  into two leaves  $(p_{\kappa_1}, \hat{y}_{\kappa_1})$  and  $(p_{\kappa_2}, \hat{y}_{\kappa_2})$  and any tree produced by further splitting  $(p_{\kappa_1}, \hat{y}_{\kappa_1})$  or  $(p_{\kappa_2}, \hat{y}_{\kappa_2})$  cannot be optimal if  $m_\kappa < 2n\lambda$ .

#### Proof.

Let  $c$  be the number of misclassifications in leaf  $(p_\kappa, \hat{y}_\kappa)$ . Since a leaf classifies according to the majority of  $m_\kappa$ , we must have

$$c \leq \frac{m_\kappa}{2} < n\lambda$$

By splitting leaf  $(p_\kappa, \hat{y}_\kappa)$  into leaves  $(p_{\kappa_1}, \hat{y}_{\kappa_1})$  and  $(p_{\kappa_2}, \hat{y}_{\kappa_2})$ , assume that we have reduced the number of misclassifications by  $b \leq c$ . Then, we have

$$\ell(f', \mathbf{X}, \mathbf{y}) = \ell(f, \mathbf{X}, \mathbf{y}) - \frac{b}{n}$$

However, we have increased the number of leaves by 1, and so

$$\lambda s(f') = \lambda s(f) + \lambda$$

Combining the last two equations, we have obtained

$$R(f', \mathbf{X}, \mathbf{y}) = R(f, \mathbf{X}, \mathbf{y}) + \lambda - \frac{b}{n}$$

However, we know that

$$\begin{aligned} b \leq c &\implies \frac{b}{n} \leq \frac{c}{n} < \frac{n\lambda}{n} = \lambda \\ &\implies -\frac{b}{n} > -\lambda \\ &\implies \lambda - \frac{b}{n} > \lambda - \lambda = 0 \end{aligned}$$

and so  $R(f', \mathbf{X}, \mathbf{y}) > R(f, \mathbf{X}, \mathbf{y})$ . This means that  $f'$  cannot be optimal according to our regularized objective. We have also proved that further splitting  $(p_{\kappa_1}, \hat{y}_{\kappa_1})$  or  $(p_{\kappa_2}, \hat{y}_{\kappa_2})$  cannot be optimal since we can just set  $f = f'$ , and apply the same argument.

### 3.1 Pruning

### 3.2 Splitting

## 4 Improved Optimization

### 4.1 GODST

## 5 Soft Decision Trees

### 5.1 Soft Splitting

### 5.2 Neural Decision Trees

## References

- [HR76] Laurent Hyafil and Ronald L. Rivest. Constructing optimal binary decision trees is np-complete. *Information Processing Letters*, 5(1):15–17, 1976.